Facilitating Research on the NVIDIA Grace Hopper Superchip: The CU Boulder Research Computing Experience

Research Computing
UNIVERSITY OF COLORADO BOULDER

Facilitating Research on the NVIDIA Grace Hopper Superchip: The CU Boulder Research Computing Experience

Presenter: Brandon Reyes

- Website: www.rc.colorado.edu
- Helpdesk: rc-help@colorado.edu
- Slides: https://github.com/ResearchComputing/rmacc_2025

Session Overview

- GH200 architecture overview
- GH200 software stack
- Beta testing phase
- Common support tasks
- Successful use cases
- Potential future directions
- How can you run on the GH200s?

GH200 architecture overview

The Grace Hopper Superchip (GH200) is a newer chip provided by NVIDIA. Its unique architecture allows the GPU and CPU to efficiently share and exchange memory.

- Ideal for workloads that utilize the CPU and GPU to a high degree and require a large amount of memory
- Although powerful, the Arm-based CPU can present support challenges
  - HPC systems often have software built only for x86_64 architectures
  - Utilizing the full capabilities can require CUDA code changes

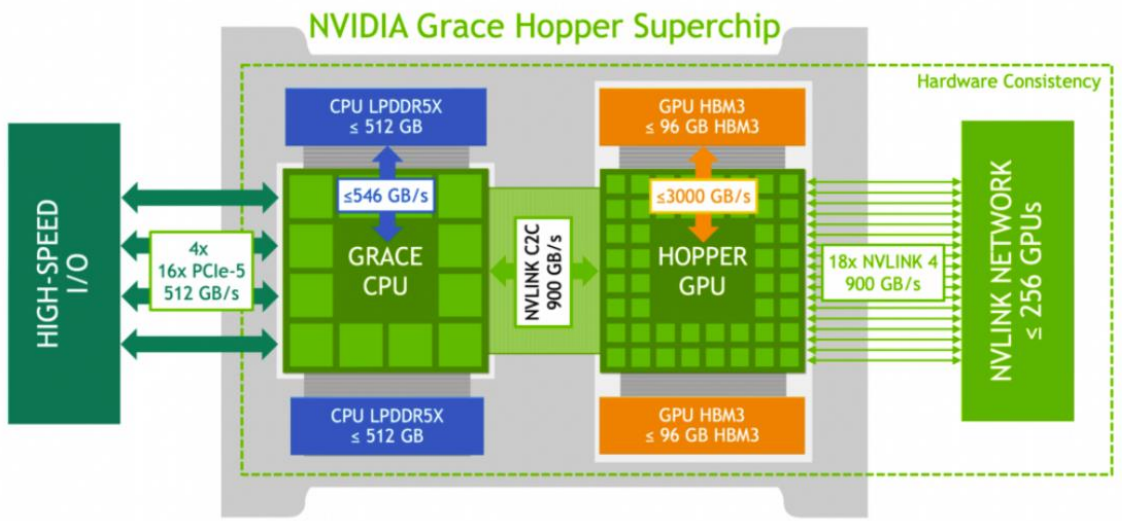

Image provided by https://resources.nvidia.com/en-us-grace-cpu/nvidia-grace-hopper

Specifications for our 2 GH200 nodes

- Grace CPU has 72 cores and roughly 480 GB of RAM
  - Arm based (Neoverse V2)
- Hopper GPU has roughly 100 GB of VRAM
- CPU and GPU are connected via NVIDIA NVLink-C2C
  - This enables efficient and seamless memory transfer between the two components
- FAST I/O speed (512 GB/s)
  - Top I/O speeds only apply to the SSD
- Roughly 1.7 TB of usable SSD on the node

GH200 software stack

Due to the unique architecture and Arm-based processors, we opted for a curated and paired-down software stack.

- CUDA compilers through NVIDIA HPC SDK e.g. nvcc, nvc++
- CUDNN libraries
- Miniforge (mamba and conda)
  - We found that most libraries that provided an aarch64 version (with GPU capabilities) worked on the GH200s
- Apptainer
  - A large selection of compatible containers are available through NVIDIA's NGC catalog

Beta testing phase

Goal: to identify any issues associated with the GH200s, software requirements, and workflows that could take advantage of the architecture

- Directed towards users with established GPU workflows
- An initial consultation was held to determine if their existing GPU workflows met the following criteria:
  - Peak GPU usage was roughly half the size of the VRAM
  - Peak CPU utilization was at least 70%
    - Exceptions were made for applications that required the large amount of available RAM
  - Software utilized was available for the GH200s

Beta testing phase cont.

Once approved, we

- Installed all necessary software and ran a trimmed-down version of the user's code
- Added users to the provided reservation
  - Enabled us to use established QoS and provide access only to approved users
- Created an allocation specific to the user
  - Facilitated easier analysis of hardware usage
- Onboarded the user to the node
  - Provided them information on how to run on the node and how to install their own software
- Provided hands-on support for any issues encountered

Common support tasks

- Creation of Mamba environments and compatible containers
- Informing users on how to move data to the local SSD
  - Most users had only used our general filesystems
- Help users understand memory consumption
  - Several users were pushing the nodes to their limit and needed to know how to monitor memory usage (e.g. “nvidia-smi”, “free -m”)
- Explaining architectural differences
  - Many users found it difficult to understand why we needed to install different software than they were already using and why some software was unavailable

Successful use cases

All users in the beta testing phase reported that they would not be able to run their workflows without the GH200s!

- All user workflows were AI based:
  - ViTs, CNNs, LSTMs, KANs, MLPs
  - Gradient boosted trees (XGBoost)
  - Image segmentation (SAM2)
- Training large models and hyperparameter tuning
  - We have seen at least 2X speedup, some users reported 10X speedup
- Inference for Large Language Models
  - We were able to run Llama 3.1 405b (requires ~300 GB of memory)

Potential future directions

- Moving out of the beta testing phase:
  - Job submissions will be limited to 1 per user
  - Node access will be provided via a QoS
  - Workflows will continue to be evaluated before permission is granted
    - Ensures proper node use and well-informed users
- We are considering utilizing Multi-Instance GPU (MIG) on one of the nodes
  - MIG would provide more resources for prototyping and testing purposes
  - We found that some users could continue to benefit from the GH200 architecture, even if they had half the GPU resources

How can you run on the GH200s?

- Submit a ticket to rc-help@colorado.edu
- In the subject provide: I am interested in running “X application” on the GH200s
- In the email body:
  - Short description of what your workflow does
  - Why you believe the GH200s would be beneficial to your workflow
  - Provide us with Linux paths to the code you would like to run
  - If possible, any JobIDs of this workflow you have ran on a GPU node

Thank you!

Slides:
https://github.com/ResearchComputing/rmacc_2025