An Approach to SLURM Configuration Verification

Kyle Reinholt

View the Slides

https://github.com/ResearchComputing/rmacc_2025/slurm_verification

Introduction – Why Configuration Matters

- Misconfigured SLURM settings often lead to user confusion, job failures, and a spike in support tickets — creating friction for both researchers and support staff.

The Problem

- Validating SLURM configs is painful
- Manual checks are error-prone
- You only get feedback *after* something breaks

The Problem

Current Approaches

- **Manual Reviews** — Time-consuming and error-prone

- **Cluster-Specific Scripts** — Custom, hard to reuse or share

- **Limited Tooling** — scontrol, slurmlint, UBCCR simulator offer partial checks

An Approach to SLURM Configuration Verification

- Declarative validation schemas (e.g. JSON schema for YAML-like configs)
- Automated testing pipelines
- Community-driven validation rules

Products Utilizing Schema Validation

Conclusion

- Config correctness is essential → current tools fall short → better tooling is possible

- Email - kyle@colorado.edu

Prior Art in SLURM Configuration Checking

- https://github.com/appeltel/slurmlint – amazing work, does not use a schema, last updated 2019

- https://ubccr-slurm-simulator.github.io/ - incredible project for simulating configurations, supports Slurm 17.11 – is currently being updated

- Spin up slurm controller in a container and see if fails

- "If it crashes, it's not right" – Joe Malingowski

Limitations with Current Approaches

- Don't scale well
- Hard to generalize
- Not integrated into CI/CD